

Class 07 Unsupervised learning

Andre Modolo

In this class we will explore clustering and dimensional reduction methods.

K-means (K=number of clusters)

Make up input data where we know what the answer should be.

```
tmp <- c(rnorm(30, -3), rnorm(30, 3))
# make it into a 2 dimensional thing
x <- cbind(x=tmp, y=rev(tmp))
# rev(tmp) flips the vector
rev(tmp)
```


[1]	2.8991538	1.6362968	3.3686467	3.6842056	1.8736334	1.9591686
[7]	3.4654623	3.1048232	2.7231510	2.3056306	2.7574076	5.8610264
[13]	3.1792991	3.8619907	3.9605646	3.5274241	4.1295692	3.4292391
[19]	2.6792261	2.1474762	4.2498852	2.7357443	1.5366285	4.0127872
[25]	1.5428144	3.7150114	2.9386080	4.9392024	2.8872510	4.6681940
[31]	-2.8379967	-1.7434181	-4.5467671	-4.2809171	-3.6596423	-4.0261784
[37]	-4.0599653	-2.9345128	-2.4536508	-4.2290156	-2.8111426	-4.1120691
[43]	-2.8259391	-4.1831423	-1.9113372	-2.8461242	-2.1304977	-4.3433768
[49]	-2.9439570	-1.9882731	-0.8469759	-1.5692192	-1.7780004	-2.8312950
[55]	-2.6384763	-2.3601613	-2.4634074	-3.8673657	-3.7020591	-2.6293157

```
head(x)
```

	x	y
[1,]	-2.629316	2.899154
[2,]	-3.702059	1.636297
[3,]	-3.867366	3.368647
[4,]	-2.463407	3.684206
[5,]	-2.360161	1.873633
[6,]	-2.638476	1.959169

Quick plot of x to see the 2 groups around (-3, 3) and (3, -3)

```
plot(x)
```



Use the `kmeans()` function setting `k` to 2 and `nstart=20` (do the picking points and finding the distances to decide a potential cluster 20 times before deciding a winning set of clusters)

```
km <- kmeans(x,center=2, nstart=20 )
km
```

K-means clustering with 2 clusters of sizes 30, 30

Cluster means:

	x	y
1	-2.985140	3.192651
2	3.192651	-2.985140

Clustering vector:

```
[1] 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 2 2 2 2 2 2
2 2
[39] 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2
```

Within cluster sum of squares by cluster:

```
[1] 58.97352 58.97352
(between_SS / total_SS = 90.7 %)
```

```
Available components:
[1] "cluster"      "centers"      "totss"        "withinss"
"tot.withinss"
[6] "betweenss"    "size"         "iter"         "ifault"
```

Clustering means: gives us the mean point of each cluster

```
km$size
```

```
[1] 30 30
```

```
km$size
```

```
[1] 30 30
```

```
[1] 30 30
```

Size of the clusters found: 30 and 30 Clustering vector: labels each component of the vector as the first or second cluster

Q. What component of your result object details? - cluster assignment/membership (1 or 2 in this case)

```
km$cluster  
[1] 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 2 2 2 2 2 2  
2 2  
[39] 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2
```

```
[1] 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 2 2 2 2 2 2  
[39] 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2
```

-Cluster center?


```
km$center
```

	x	y
1	-2.985140	3.192651
2	3.192651	-2.985140

	x	y
1	-2.985140	3.192651
2	3.192651	-2.985140

Q. plot x colored by the kmeans cluster assignment and add cluster centers as blue points

```
plot(x, col=km$cluster)
points(km$centers, col="blue", pch=15, cex=2)
```



What if I ask for more than 2 clusters?

```
km4 <- kmeans(x, 4, nstart=20)
km4
```

K-means clustering with 4 clusters of sizes 12, 12, 18, 18

Cluster means:

	x	y
1	-3.825862	3.975168
2	3.975168	-3.825862
3	-2.424659	2.670972
4	2.670972	-2.424659

Clustering vector:

```
[1] 3 3 1 3 3 3 3 3 3 3 1 1 3 1 3 1 3 1 3 3 1 3 1 1 1 3 1 2 4 2 2 2 4
2 4
[39] 4 2 4 2 4 2 4 2 4 2 2 4 4 4 4 4 4 4 4 2 4 4
```


Within cluster sum of squares by cluster:

```
[1] 13.05430 13.05430 19.53628 19.53628
(between_SS / total_SS = 94.8 %)
```

Available components:

```
[1] "cluster"      "centers"      "totss"        "withinss"
"tot.withinss"
[6] "betweenss"    "size"         "iter"         "ifault"

plot(x, col=km4$cluster)
points(km4$centers, col="blue", pch=15, cex=2)
```



#Hierarchical Clustering

Super useful and widely employed clustering method which has the advantage over kmeans because it can show you a little something about the true nature of the clustering in your data. You need to give it a “d” distance matrix as an input (how far apart the values are). get it using `dist()`


```
d <- dist(x)
hc <- hclust(d)
hc
```

Call:
`hclust(d = d)`

```
Cluster method   : complete
Distance         : euclidean
Number of objects: 60
```

There is a plot method for hcluster results:

```
plot(hc)
```




You get 2 overall branches with 1:30 on one branch and 31:60 on the other branch. This makes sense because in the vector we made the first 30 numbers have a set mean and the second 30 numbers have another set mean.

Long goal post = big jump between the things you grouped together and the next group.

How do I get an actual result out of this? cut the longest post, and you are left with “subtrees” in this case you are left with 2 subtrees.

```
plot(hc)
#cut the tree with this line
abline(h=10, col="red")
```

Principal Component Analysis (PCA)

The base R function to do PCA is called `prcomp()`

Import the food data from the 4 countries

```
url <- "https://tinyurl.com/UK-foods"
x <- read.csv(url)
dim(x)
```

```
[1] 17  5
```

There are 17 row and 5 columns

```
head(x)
```

	X	England	Wales	Scotland	N.Ireland
1	Cheese	105	103	103	66
2	Carcass_meat	245	227	242	267
3	Other_meat	685	803	750	586
4	Fish	147	160	122	93
5	Fats_and_oils	193	235	184	209
6	Sugars	156	175	147	139

We can remove the x column and only get the 4 counties as columns by using this code

```
rownames(x) <- x[,1]  
x<- x[,-1]  
head(x)
```

	England	Wales	Scotland	N.Ireland
Cheese	105	103	103	66
Carcass_meat	245	227	242	267
Other_meat	685	803	750	586
Fish	147	160	122	93
Fats_and_oils	193	235	184	209
Sugars	156	175	147	139

Be careful with this approach because if you keep running it multiple times, it will keep removing the name of the first column


A more robust way of doing it would be using this method, setting the row name as 1, so you can rerun the code and it won't delete any more column names

```
x <- read.csv(url, row.names=1)  
head(x)
```

	England	Wales	Scotland	N.Ireland
Cheese	105	103	103	66
Carcass_meat	245	227	242	267
Other_meat	685	803	750	586
Fish	147	160	122	93
Fats_and_oils	193	235	184	209
Sugars	156	175	147	139


Spotting the major differences and trends using a bar plot

```
barplot(as.matrix(x), beside=T, col=rainbow(nrow(x)))
```




Doing `beside=False` you get this kind of bar plot

```
barplot(as.matrix(x), beside=F, col=rainbow(nrow(x)))
```



What about plotting it this way?

```
pairs(x, col=rainbow(10), pch=16)
```



What does it mean when a point lies on a diagonal of a given plot? This gives a matrix of scatterplots comparing the countries as an x and a y variable in each situation. This way you only have to look at bottom left or top right half depending in which country you want to be on which axis.

If the point lies on the diagonal of a scatterplot, this means that the two countries have a similar amount of consumption for that specific food group (color)

The main difference in food consumption between N. Ireland and the other countries is in the food colored blue

#PCA to the rescue

Take the transpose of x to flip the rows and columns

```
t(x)
```

	Cheese	Carcass_meat	Other_meat	Fish	Fats_and_oils	Sugars
England	105	245	685	147	193	156
Wales	103	227	803	160	235	175
Scotland	103	242	750	122	184	147
N.Ireland	66	267	586	93	209	139

	Fresh_potatoes	Fresh_Veg	Other_Veg	Processed_potatoes
England	720	253	488	198

Wales	874	265	570		203
Scotland	566	171	418		220
N.Ireland	1033	143	355		187
	Processed_Veg	Fresh_fruit	Cereals	Beverages	Soft_drinks
England	360	1102	1472	57	1374
Wales	365	1137	1582	73	1256
Scotland	337	957	1462	53	1572
N.Ireland	334	674	1494	47	1506
	Alcoholic_drinks	Confectionery			
England	375	54			
Wales	475	64			
Scotland	458	62			
N.Ireland	135	41			

Now do `prcomp()` and print out the summary

```
pca <- prcomp(t(x))
summary(pca)
```

Importance of components:

	PC1	PC2	PC3	PC4
Standard deviation	324.1502	212.7478	73.87622	4.189e-14
Proportion of Variance	0.6744	0.2905	0.03503	0.000e+00
Cumulative Proportion	0.6744	0.9650	1.00000	1.000e+00

Proportion of Variance: 67.4% of all the variance is captured on the new axis made.

Cumulative Proportion: adding 2 or 3 PCs together you capture basically all the variance from the plots (ex: PC2 with 96.5%!)


A “PCA plot” (a.k.a “Score Plot”, PC1vsPC2 plot, etc.)

```
pca$x
```

	PC1	PC2	PC3	PC4
England	-144.99315	2.532999	-105.768945	2.842865e-14
Wales	-240.52915	224.646925	56.475555	7.804382e-13
Scotland	-91.86934	-286.081786	44.415495	-9.614462e-13
N.Ireland	477.39164	58.901862	4.877895	1.448078e-13

Plot the PC1 vs PC2 and color the countries Ireland is green

```
plot(pca$x[,1], pca$x[,2], col=c("orange", "red", "blue", "green"), pch=18)
```




You see that N. Ireland is actually different than the other countries in their food consumption.

Below we can use the square of `pca$sdev`, which stands for “standard deviation”, to calculate how much variation in the original data each PC accounts for:

```
v <- round( pca$sdev^2/sum(pca$sdev^2) * 100 )
v
[1] 67 29 4 0
```

This information can be summarized in a plot of the variances (eigenvalues) with respect to the principal component number (eigenvector number), which is given below.


```
barplot(v, xlab="Principal Component", ylab="Percent Variation")
```



We can also consider the influence of each of the original variables upon the principal components (typically known as loading scores). This information can be obtained from the `prcomp()` returned `$rotation` component

Using PC1 we can get this barplot:


```
par(mar=c(10, 3, 0.35, 0))  
barplot( pca$rotation[,1], las=2 )
```



Now we can see what foods that make N. Ireland more different than the rest of the countries.

Generate a similar 'loadings plot' for PC2. What two food groups feature prominently and what does PC2 mainly tell us about?

```
par(mar=c(10, 3, 0.35, 0))
barplot( pca$rotation[,2], las=2 )
```

Using PC2 it tells us that N. Ireland eats more fresh potatoes and drinks less soft drinks